

# Adversarial Machine Learning

Machine Learning Course - CS-433

Nov 13, 2024

Nicolas Flammarion

**EPFL**



# Some input examples are hard for humans



Dog or mop?

- Some examples may be challenging for humans
- NNs typically have **no problem** with them
- However, NNs are not always **robust** in their decisions

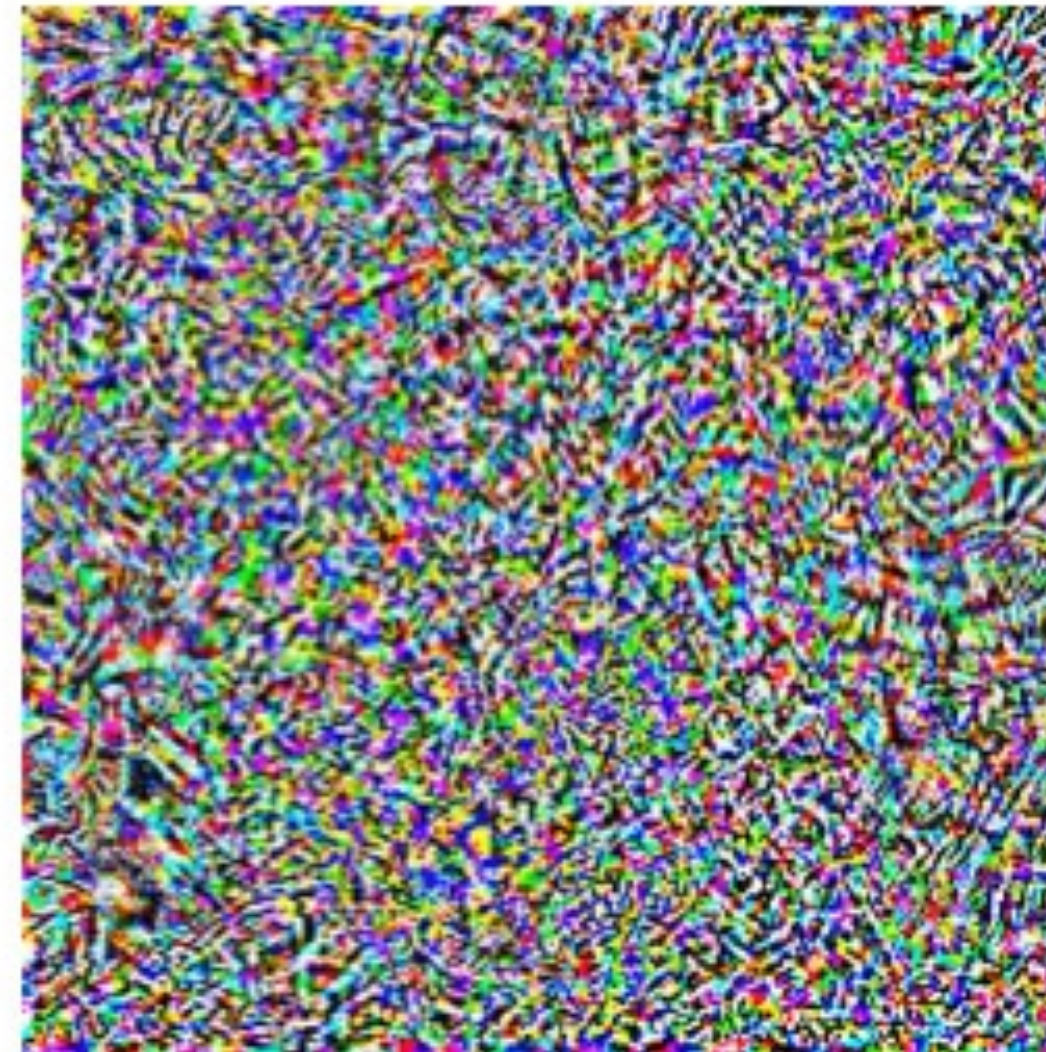


# Adversarial examples are small perturbations that cause misclassification with high confidence

“pig”



+ 0.005 x



=



“airliner”

Source: Z. Kolter, A. Madry, NeurIPS'18 tutorial on adversarial robustness

NNs have struggle with imperceptible yet very specific inputs known as **adversarial examples**


- ➡ **Security issue:** consider the implications for a self-driving car and its ability to detect stop signs
- ➡ We don't understand how these models **generalize** and react to shifts in the distribution of data (i.e., **distribution shifts**)

# Standard risk

Classification problem:  $(X, Y) \sim \mathcal{D}$ ,  $Y$  with range  $\{-1, 1\}$

Standard risk: average zero-one loss over  $X$

$$R(f) = \mathbb{E}_{\mathcal{D}} \left[ 1_{f(X) \neq Y} \right] = \mathbb{P}_{\mathcal{D}} \left[ f(X) \neq Y \right]$$





# Standard risk vs. adversarial risk

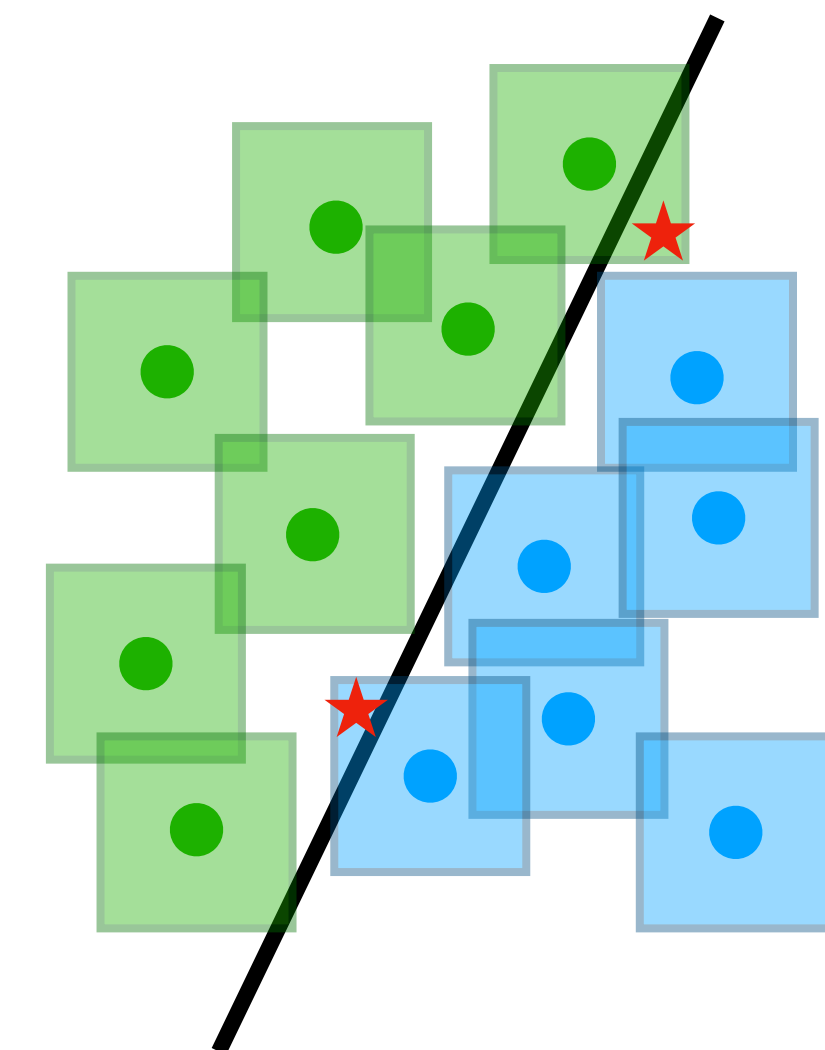
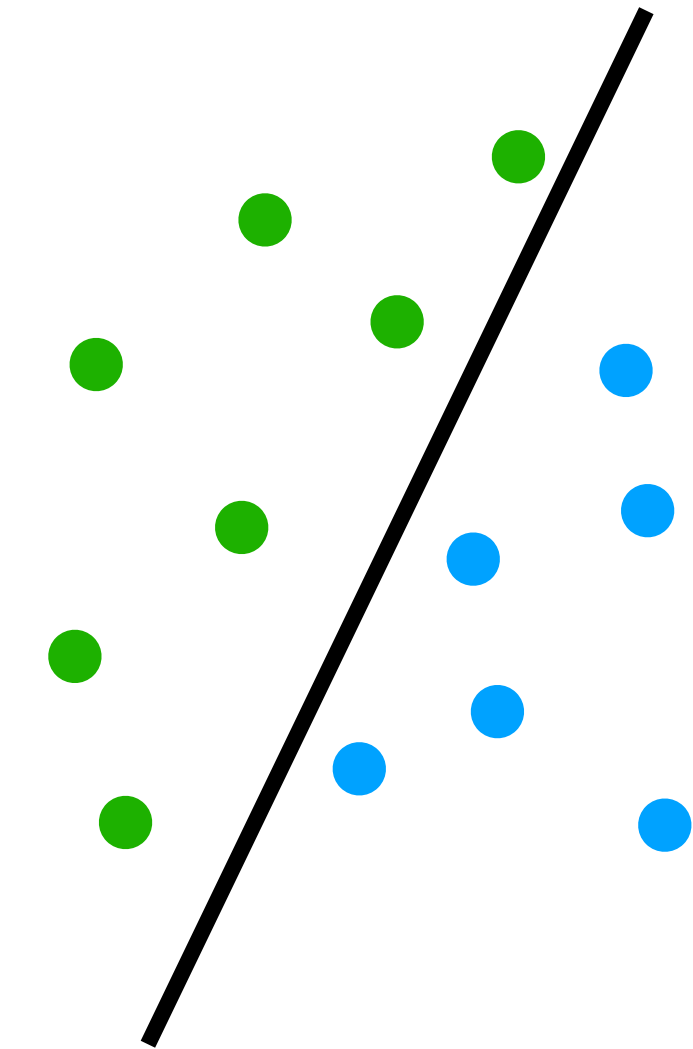
Classification problem:  $(X, Y) \sim \mathcal{D}$ ,  $Y$  with range  $\{-1, 1\}$

Standard risk: average zero-one loss over  $X$

$$R(f) = \mathbb{E}_{\mathcal{D}} \left[ 1_{f(X) \neq Y} \right] = \mathbb{P}_{\mathcal{D}} [f(X) \neq Y]$$

Adversarial risk: average zero-one loss over **small, worst-case perturbations of  $X$**

$$R_{\varepsilon}(f) = \mathbb{E}_{\mathcal{D}} \left[ \max_{\hat{x}, \|\hat{x} - X\| \leq \varepsilon} 1_{f(\hat{x}) \neq Y} \right]$$



# Adversarial vulnerability raises many questions

$$R_\varepsilon(f) = \mathbb{E}_{\mathcal{D}} \left[ \max_{\hat{x}, \|\hat{x}-X\| \leq \varepsilon} 1_{f(\hat{x}) \neq Y} \right]$$

- Threat model:
  - How should we define the adversary's power?
  - Which norm should we consider?  $\ell_\infty$ ,  $\ell_2$ ,  $\ell_1$ ,  $\ell_0$ , ...
  - What set of perturbations?
- If  $R(f) \leq \delta$ , then how large can  $R_\varepsilon(f)$  be?

# Adversarial vulnerability raises many questions

- How can we compute an adversarial example?
- What level of access do we have to the model to attack it?
- How can we design a classifier  $f$  so that it is robust?  
Related: given a non-robust classifier, how can we make it robust?
- Why are neural networks non-robust?

# Generating adversarial examples

Task: given an input  $(x, y)$  and a model  $f: \mathcal{X} \rightarrow \{-1, 1\}$ , find an input  $\hat{x}$ , such that

(a)  $\|\hat{x} - x\| \leq \varepsilon$

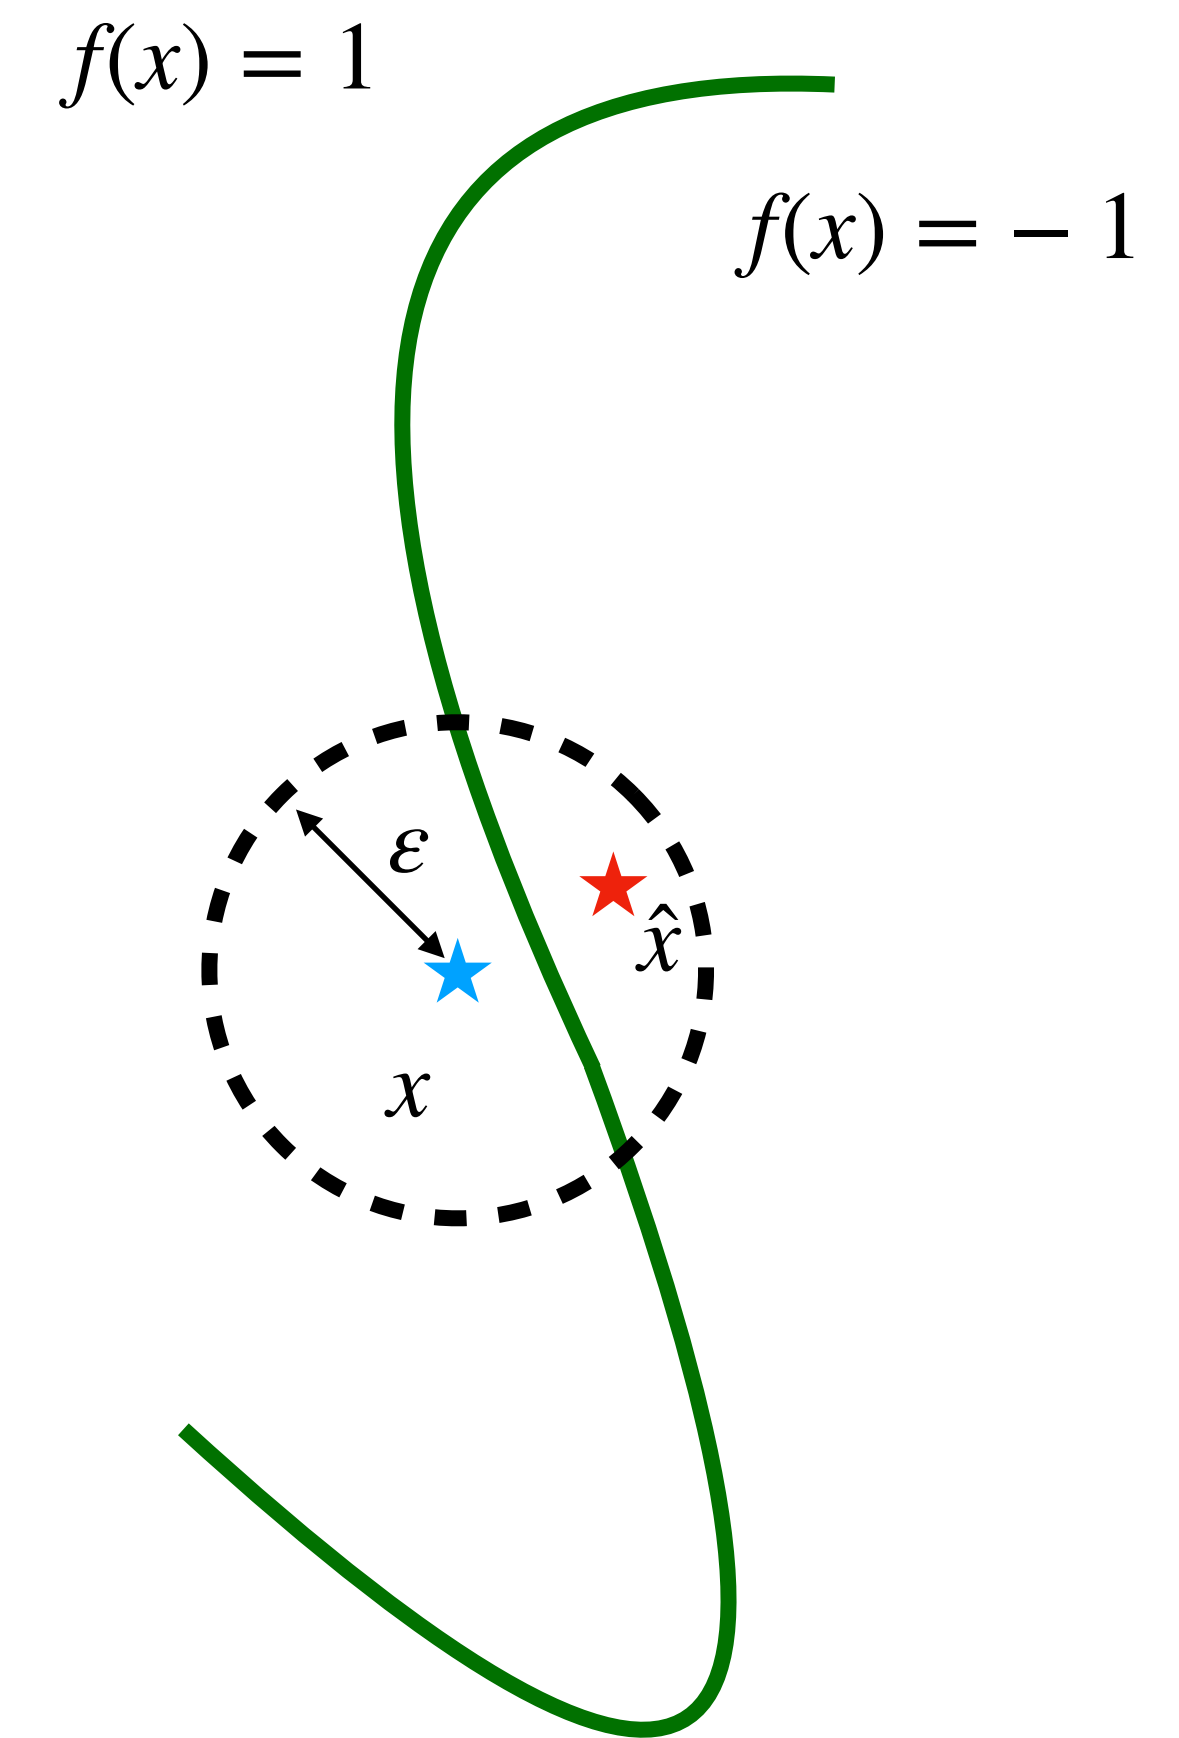
(b) the model  $f$  makes a mistake on it

**Trivial case**:  $x$  is already misclassified

➡ No action required

**General case**:  $x$  is correctly classified

➡ find  $\hat{x}$  such that  $f(\hat{x}) \neq y$  and  $\|\hat{x} - x\| \leq \varepsilon$   
i.e.,  $\hat{x} \in B_x(\varepsilon) \cap \{x', f(x') = -y\}$





# Generating adversarial examples amounts to maximizing the classification loss w.r.t. the inputs

Find an adversarial example by solving

$$\max_{\hat{x}, \|\hat{x}-x\|\leq\epsilon} 1_{f(\hat{x})\neq y}$$

➡ Optimization problem with respect to the inputs

Problem: optimizing the indicator function  $1_{f(\hat{x})\neq y}$  is difficult:

1. The indicator function  $1$  is not continuous
2. The NN prediction  $f$  outputs discrete class values  $\{-1, 1\}$


# Generating adversarial examples amounts to solving a constrained optimization problem

## Solution:

1. Use a smooth classification loss  $\ell$  (e.g., logistic or hinge loss) instead
2. Consider the output  $g$  of the NN before classification (i.e.,  $f(x) = \text{sign}(g(x))$ )

**Main idea:** Replace the difficult problem involving the indicator with a smooth problem

$$\max_{\hat{x}, \|\hat{x}-x\| \leq \varepsilon} \mathbb{1}_{f(\hat{x}) \neq y} \longrightarrow \max_{\hat{x}, \|\hat{x}-x\| \leq \varepsilon} \ell(yg(\hat{x}))$$



**Reminder:** decreasing, margin-based (i.e., dependent on  $y \cdot g(x)$ ) classification losses

# Generating adversarial examples: white-box case


How to solve  $\max_{\hat{x}, \|\hat{x}-x\| \leq \epsilon} \ell(yg(\hat{x}))$  in the **white-box** case, i.e., if we know the model  $g$ ?

Compute its gradient:  $\nabla_x \ell(yg(x)) = \underbrace{y\ell'(yg(x))}_{\leq 0 \text{ since classification loss are decreasing}} \nabla_x g(x)$

We should move in the direction  $\propto -y \nabla_x g(x)$

Interpretation:  $f(x) = \text{sign}(g(x))$

- If  $y = 1$ , we want to decrease  $g(x)$  and follow  $-\nabla_x g(x)$
- If  $y = -1$ , we want to increase  $g(x)$  and follow  $\nabla_x g(x)$

 Why use  $\ell$ , and not directly minimize  $yg(\hat{x})$ ?  
→ It won't extend to multi-class classification and robust training.



# Generating adversarial examples: taking into account the constraints

We can linearize the loss  $\tilde{\ell}(x) := \ell(yg(x))$  to derive an iteration:

$$\begin{aligned}\max_{\|\hat{x}-x\|\leq\epsilon} \tilde{\ell}(\hat{x}) &\approx \max_{\|\hat{x}-x\|\leq\epsilon} \tilde{\ell}(x) + \nabla_x \tilde{\ell}(x)^T (\hat{x} - x) \\ &= \tilde{\ell}(x) + \max_{\|\hat{x}-x\|\leq\epsilon} \nabla_x \tilde{\ell}(x)^T (\hat{x} - x) \\ &= \tilde{\ell}(x) + \max_{\|\delta\|\leq\epsilon} \nabla_x \tilde{\ell}(x)^T \delta\end{aligned}$$

- We need to maximize the inner product under a norm constraint, i.e. find the optimal local update
- This is a simple problem for which we can get a closed-form solution depending on the norm used to measure the perturbation size  $\|\delta\|$

# Generating adversarial examples: one-step attack

## Problem:

$$\max_{\|\delta\| \leq \varepsilon} \nabla_x \tilde{\ell}(x)^T \delta$$



- Solution for the  $\ell_2$  norm:  $\delta_2^* = \varepsilon \cdot \frac{\nabla_x \tilde{\ell}(x)}{\|\nabla_x \tilde{\ell}(x)\|_2} = -\varepsilon y \cdot \frac{\nabla_x g(x)}{\|\nabla_x g(x)\|_2}$

$$\Rightarrow \hat{x} = x - \varepsilon y \cdot \frac{\nabla_x g(x)}{\|\nabla_x g(x)\|_2}$$

- Solution for the  $\ell_\infty$  norm:  $\delta_\infty^* = \varepsilon \cdot \text{sign}(\nabla_x \tilde{\ell}(x)) = -\varepsilon y \cdot \text{sign}(\nabla_x g(x))$

$$\Rightarrow \hat{x} = x - \varepsilon y \cdot \text{sign}(\nabla_x g(x))$$

- **Fast Gradient Sign Method**  
[Goodfellow et al., 2014]



# Generating adversarial examples: multi-step attack

These updates can be done iteratively and combined with a projection  $\Pi$  on the feasible set (i.e.,  $\ell_2/\ell_\infty$  balls here)

## Projected Gradient Descent (PGD attack):

- $\ell_2$  norm:



$$\delta^{t+1} = \Pi_{B_2(\varepsilon)} \left[ \delta^t + \alpha \cdot \frac{\nabla \tilde{\ell}(x + \delta^t)}{\|\nabla \tilde{\ell}(x + \delta^t)\|_2} \right],$$

$$\text{where } \Pi_{B_2(\varepsilon)}(\delta) = \begin{cases} \varepsilon \cdot \delta / \|\delta\|_2, & \text{if } \|\delta\|_2 \geq \varepsilon \\ \delta, & \text{otherwise} \end{cases}$$

- $\ell_\infty$  norm:

$$\delta^{t+1} = \Pi_{B_\infty(\varepsilon)} \left[ \delta^t + \alpha \cdot \text{sign}(\nabla \tilde{\ell}(x + \delta^t)) \right],$$


$$\text{where } \Pi_{B_\infty(\varepsilon)}(\delta)_i = \begin{cases} \varepsilon \cdot \text{sign}(\delta_i), & \text{if } |\delta_i| \geq \varepsilon \\ \delta_i, & \text{otherwise} \end{cases}$$





# Reminder: $\ell_p$ norms


Different  $\ell_p$  norms have different geometry



The difference is especially pronounced in high dimensions!

# Visualizations of different $\ell_p$ adversarial examples

The choice of the norm leads to different properties of the resulting adversarial perturbations: e.g.  $\ell_\infty$  are **dense** and  $\ell_0$  are **sparse**



Source: Towards Evaluating the Robustness of Neural Networks, Carlini et al., 2018

Which perturbations should we aim to be robust against?

➡ Extensive research on formulating the 'right' perturbation set!

# White-box attacks: implementation

- For a neural network, the gradients  $\nabla_x g(x)$  can also be computed by **backpropagation** (note: they are taken w.r.t. **inputs**, not parameters!)
- Modern deep learning frameworks readily support this  
→ **lab #9** (implement Fast Gradient Sign Method on MNIST in PyTorch)
- Now: what **if we don't know**  $g(x)$ ? Is it possible to run an attack without knowing how to compute the gradient  $\nabla_x g(x)$ ?



# Black-box attacks: query-based gradient estimation

There are different assumptions on the knowledge about the model  $f$ :

- **score-based**: we can query the continuous model scores  $g(x) \in \mathbb{R}$
- **decision-based**: we can query only the predicted class  $f(x) \in \{-1, 1\}$

In the score-based case, we can approximate the gradient by using the finite difference formula:

$$\nabla_x g(x) \approx \sum_{i=1}^d \frac{g(x + \alpha e_i) - g(x)}{\alpha} e_i$$

Remark: similar techniques can be adapted for the decision-based case (when  $x$  is near the decision boundary)

# Black-box attacks via transfer attacks

Alternative approach: **transfer attacks**

1. train a **similar** surrogate model  $\hat{f} \approx f$  on **similar** data
  2. transfer the resulting white-box adversarial perturbation from  $\hat{f}$  to  $f$
- Success depends on how **similar** the model architecture and data are
  - If we are allowed to query  $f$  given some **unlabeled** inputs  $\{x_n\}_{n=1}^N$ , we can obtain  $\{x_n, f(x_n)\}_{n=1}^N$  and use that information to learn  $\hat{f}$  (known as **model stealing**)
    - can facilitate **transfer attacks**

# Black-box attacks: summary

**General takeaway:** black-box attacks are of practical concern but:

- Query-based methods often require a lot of queries (10k-100k), particularly **decision-based** attacks → easy to restrict access for the attacker!
- Obtaining a surrogate model  $\hat{f}$  can be costly and there is no guarantee of success
- A critical missing element is the implementation of **physically realizable attacks**



# Physically realizable attacks

For practical application, adversarial examples must meet additional requirements:

- Resilience to JPEG compression (for images input directly in a digital format)
- Resilience to photographic distortions (for real-world adversarial examples captured by a camera)
- Resilience to varying camera angles (for a moving camera, e.g., on a self-driving car)

→ a surge of papers on how to take these requirements into account



Source: Robust Physical-World Attacks on Deep Learning Visual Classification (CVPR 2018)

# How do we train robust models?

We have seen how to **generate** adversarial examples, but **how can we make our models robust to such attacks?**

➡ Simply train them on these adversarial examples, a.k.a. **adversarial training**

- **Standard training:** the goal is to minimize the **standard risk**:

$$\min_{\theta} R(f_{\theta}) = \mathbb{E}_{\mathcal{D}} \left[ 1_{f(X) \neq Y} \right]$$

- **Adversarial training:** the goal is to minimize the **adversarial risk**:

$$\min_{\theta} R_{\varepsilon}(f_{\theta}) = \mathbb{E}_{\mathcal{D}} \left[ \max_{\hat{x}, \|\hat{x} - X\| \leq \varepsilon} 1_{f(\hat{x}) \neq Y} \right]$$

# Adversarial training: formulation

**Goal:**

$$\min_{\theta} R_{\varepsilon}(f_{\theta}) = \mathbb{E}_{\mathcal{D}} \left[ \max_{\hat{x}, \|\hat{x}-X\| \leq \varepsilon} 1_{f(\hat{x}) \neq Y} \right]$$

- The data distribution  $\mathcal{D}$  is unknown  $\rightarrow$  approximate it with a **sample average**
- The classification loss is non-continuous  $\rightarrow$  use a **smooth loss**

This leads to the following **robust optimization** problem:


$$\min_{\theta} \frac{1}{N} \sum_{n=1}^N \max_{\hat{x}_n, \|x_n - \hat{x}_n\| \leq \varepsilon} \ell(y_n g_{\theta}(\hat{x}_n))$$

**Interpretation:** minimize the risk on adversarial examples

# Adversarial training: algorithm

$$\min_{\theta} \frac{1}{N} \sum_{n=1}^N \max_{\hat{x}_n, \|x_n - \hat{x}_n\| \leq \varepsilon} \ell(y_n g_{\theta}(\hat{x}_n))$$

**Adversarial training:** at each iteration  $t$ :


1. For each  $x_n$ , approximate  $\hat{x}_n^{\star} \approx \arg \max_{\|x_n - \hat{x}_n\| \leq \varepsilon} \ell(y_n g_{\theta}(\hat{x}_n))$  via the **PGD attack**
2. Perform a gradient descent step w.r.t.  $\theta$  using  $\frac{1}{N} \sum_{n=1}^N \nabla_{\theta} \ell(y_n g_{\theta}(\hat{x}_n^{\star}))$   
 Note you are using  $\hat{x}_n^{\star}$  and not  $x_n$



# Adversarial training: discussion

## Good news:


- Adversarial training is a state-of-the-art approach for robust classification!
- Adversarial training leads to **more interpretable gradients**  $\nabla_x \ell(yg_\theta(x))$
- The algorithm is fully compatible with SGD  
→ you will explore it in **lab #9**  
(adversarial training of a CNN on MNIST)



# Adversarial training: discussion

## Bad news:

- Increased computational time: proportional to the number of PGD steps
- **Robustness-accuracy tradeoff:** using too large  $\epsilon$  lead to worse standard accuracy (right)



Source: Robustness May Be at Odds with Accuracy (ICLR 2019)

# Key question: so why do adversarial examples exist?



We can conceptualize it with a simple model

Consider  $x \in \mathbb{R}^d$ ,  $y \sim \mathcal{U}(\{-1, 1\})$ ,  $Z_i \sim \mathcal{N}(0, 1)$ :


- **Robust features:**  $x_1 = y + Z_1$
- **Non-robust features:**  $x_i = y\sqrt{\frac{\log d}{d-1}} + Z_i$  for  $i \in \{2, \dots, d\}$

We'll see that when  $d \rightarrow \infty$ :

- **standard risk** can be arbitrarily **small**
- **adversarial risk** can be arbitrarily **large**



# Model is only using the robust feature $x_1$



assuming  $p(y = 1) = p(y = -1)$

**MLE:**  $\arg \max_{\hat{y} \in \{\pm 1\}} p(\hat{y} \mid x_1) = \arg \max_{\hat{y} \in \{\pm 1\}} \frac{p(x_1 \mid \hat{y})p(\hat{y})}{p(x_1)} = \arg \max_{\hat{y} \in \{\pm 1\}} p(x_1 \mid \hat{y})$

**Standard risk:**  $\int_0^\infty \frac{1}{\sqrt{2\pi}} e^{-0.5(x+1)^2} dx \approx 0.16 \rightarrow \text{good but not perfect!}$

# Model is using both robust and non-robust features (I)

Let's derive MLE using **all** features using the shortcut notation  $x_i = ya_i + Z_i$ :

$$\begin{aligned}\arg \max_{\hat{y} \in \{\pm 1\}} p(\hat{y} \mid x) &= \arg \max_{\hat{y} \in \{\pm 1\}} \prod_{i=1}^d p(x_i \mid \hat{y}) \\&= \arg \max_{\hat{y} \in \{\pm 1\}} \sum_{i=1}^d \log p(x_i \mid \hat{y}) \\&= \arg \max_{\hat{y} \in \{\pm 1\}} \sum_{i=1}^d \log \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}(x_i - \hat{y}a_i)^2} \\&= \arg \min_{\hat{y} \in \{\pm 1\}} \sum_{i=1}^d (x_i - \hat{y}a_i)^2 \\&= \arg \min_{\hat{y} \in \{\pm 1\}} \sum_{i=1}^d (x_i^2 - 2x_i\hat{y}a_i + \hat{y}^2a_i^2) = \arg \max_{\hat{y} \in \{\pm 1\}} \hat{y} \sum_{i=1}^d x_i a_i\end{aligned}$$



# Model is using both robust and non-robust features (II)

The MLE expression we maximize over  $\hat{y} \in \{-1, 1\}$  becomes:

$$\hat{y} \sum_{i=1}^d x_i a_i = \hat{y} y \left( \sum_{i=1}^d a_i^2 \right) + \hat{y} \sum_{i=1}^d a_i Z_i = \hat{y} y (1 + \log(d)) + \hat{y} Z,$$

where  $Z := \sum_{i=1}^d a_i Z_i \sim \mathcal{N}(0, 1 + \log d)$

Scaling by  $1/(1 + \log d)$ , the MLE expression results in:

$$y\hat{y} + \hat{y}Z \text{ with } Z \sim \mathcal{N}(0, 1/(1 + \log d))$$

**Conclusion:** when the dimension  $d \rightarrow \infty$ ,  $\hat{y}Z \rightarrow 0$  and standard risk  $\rightarrow 0$

**Interpretation:** using the non-robust features improves standard risk!

# What about adversarial risk?

- The adversary can use tiny  $\ell_\infty$ -perturbations

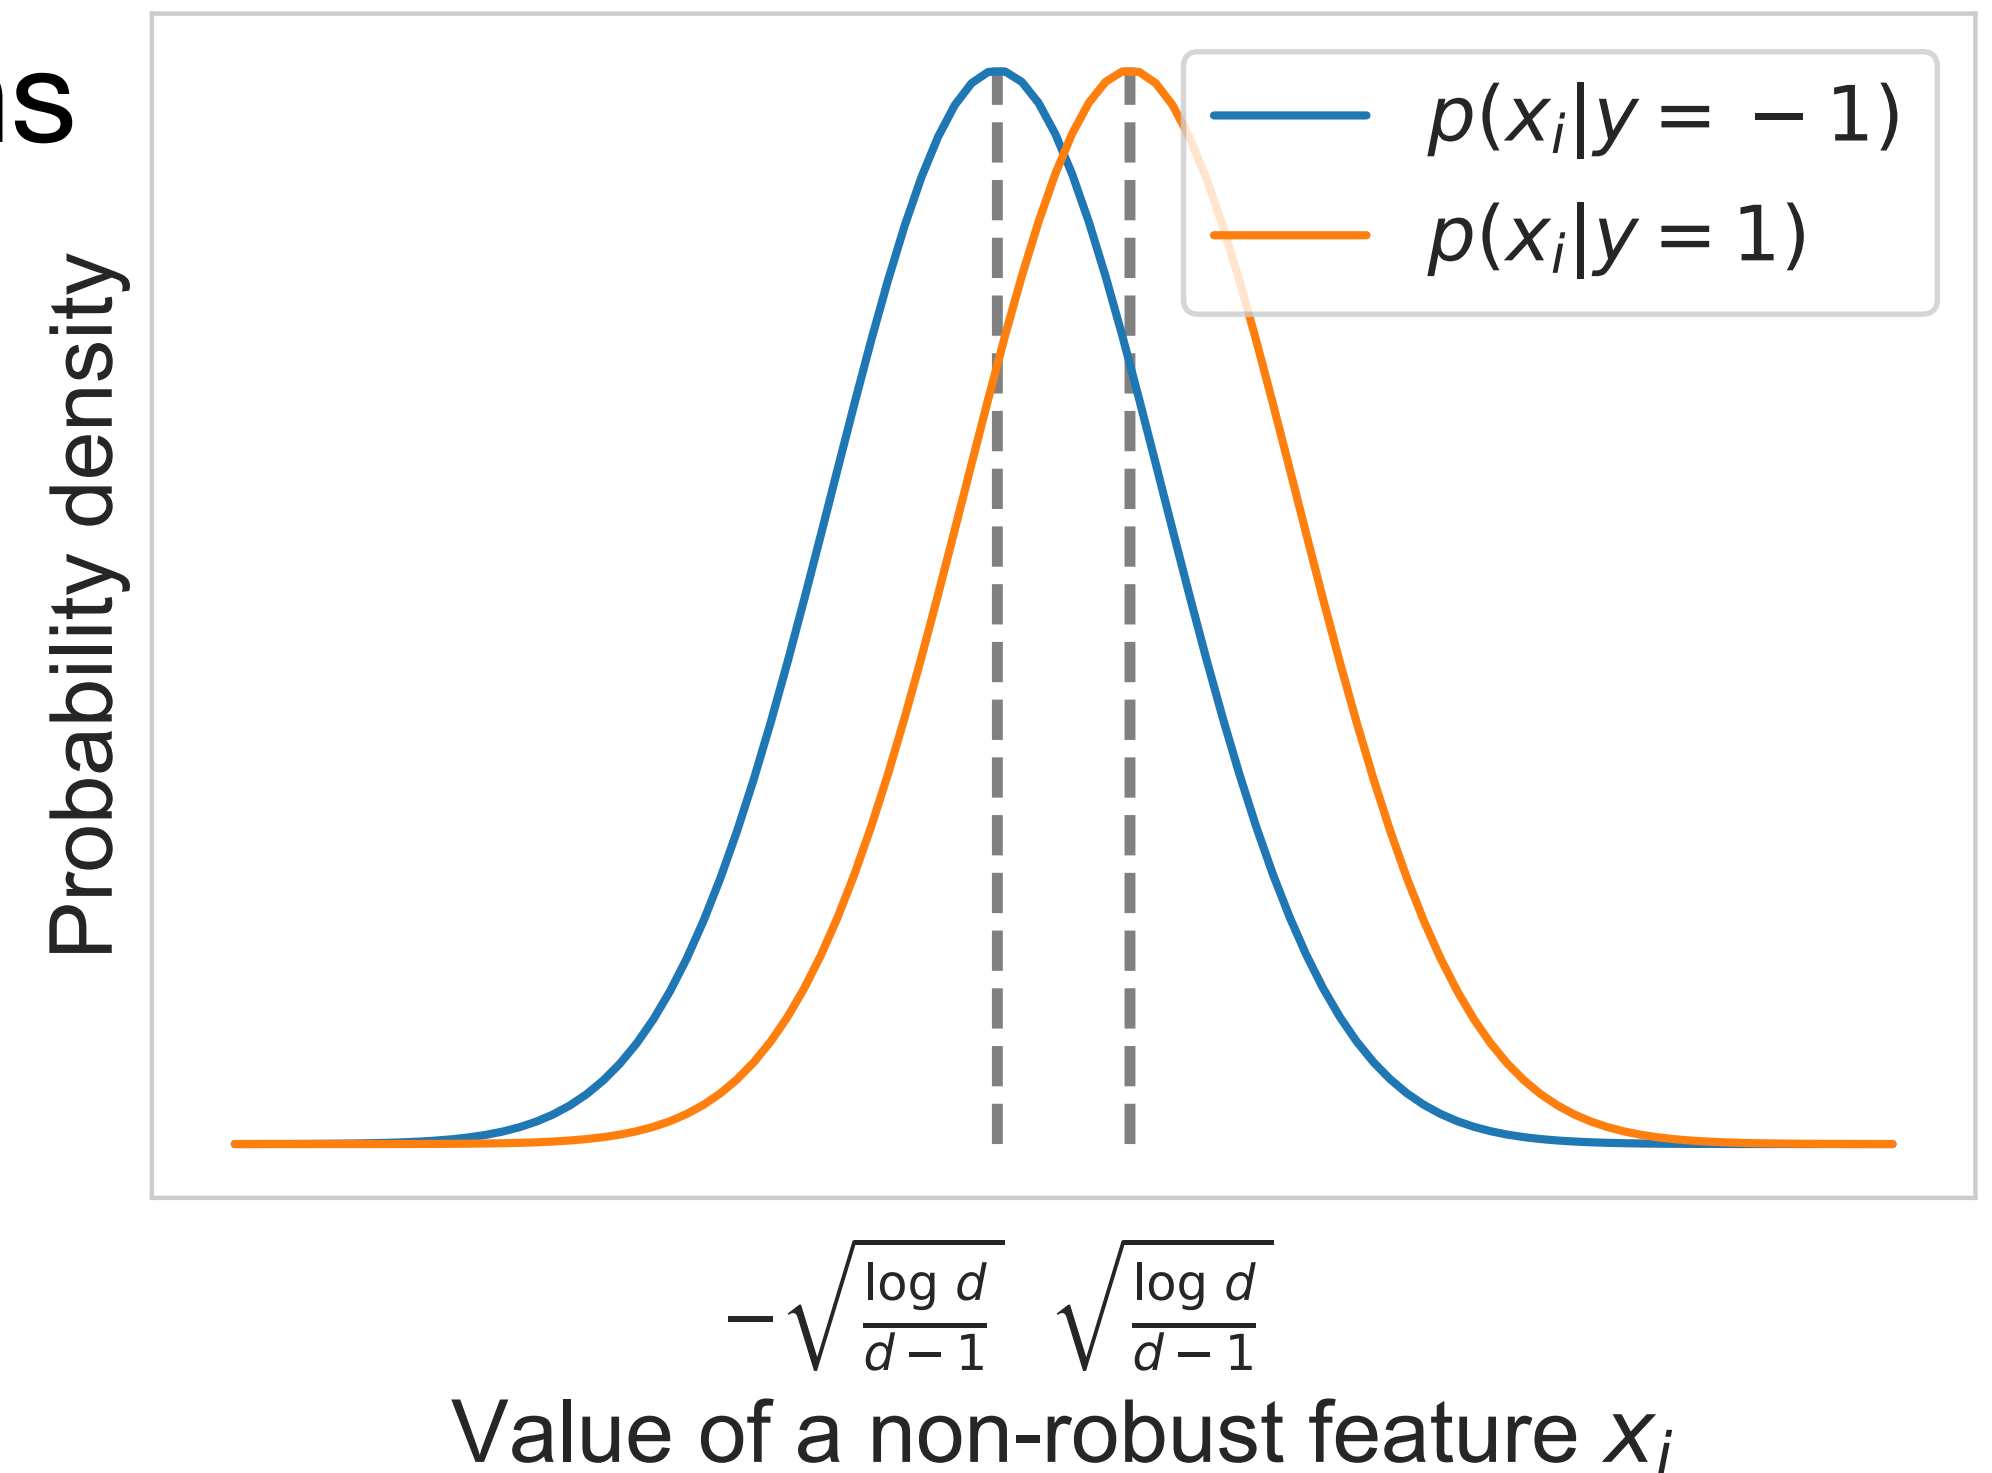
$$\varepsilon = 2\sqrt{\frac{\log d}{d-1}} \quad (\rightarrow 0 \text{ when } d \rightarrow \infty)$$

- Optimal adversarial strategy:

$$\hat{x}_1 = \left(1 - 2\sqrt{\frac{\log d}{d-1}}\right)y + Z_1 \text{ (almost unaffected)}$$

$$\hat{x}_i = -\sqrt{\frac{\log d}{d-1}}y + Z_i \text{ (completely flipped!)}$$

- Adversarial risk**  $R_\varepsilon(f)$  will become  $\approx 1$   
(due to non-robust  $x_i$ ) although **standard risk**  $R(f)$  is 0!
- But:** only using the robust feature  $x_1$  leads to  $R_\varepsilon(f) \approx R(f) = 0.16$   
➡ **tradeoff** between accuracy and robustness



# Recap

- **NNs may be susceptible to adversarial examples** imperceptible to us
- Adversarial examples can be **obtained via gradient steps on the input**
- Adversarial training: Enhance model robustness by **training on adversarially perturbed input data**
- Robustness typically comes at the cost of accuracy, since **non-robust features can still be useful features**